The All-Star Equation: Decoding Basketball's Elite

Daniel Ma, Jeffrey Liang, Augustus Bailey, Jaden Chung

Team: BALLHOGS

Introduction

What really sets NBA All-Stars apart? We explored the **statistical distinctions between NBA All-Stars and average players**, investigating **which performance metrics truly separate the elite from the rest of the league**. Through rigorous analysis of comprehensive player statistics across 20+ seasons, we aim to identify the key performance indicators that best predict All-Star selection, providing insights into the quantifiable aspects of basketball excellence that transcend traditional evaluation methods.

Hypothesis

Metrics that suggest heavy offensive involvement (such as PTS and Usage Rate) will be strongly correlated with All-Star Status, especially for players who have team success (such as high Net Rating). Additionally, these usage metrics are strong predictors for future NBA All-Stars.

Bradley Beal (2019-20)
- Pre-All-Star PPG: **29.1**
- Pre-All-Star Net Rating: **-4.9**
- All-Star?: **NO**

Stephen Curry (2015-16)
- Pre-All-Star PPG: **29.8**
- Pre-All-Star Net Rating: **+21.0**
- All-Star?: **YES**

Methodology & Data Collection

- Dataset consists of **over 11,000 rows** gathered from the **official NBA stats API**.
- Collected data for every NBA player during the **pre-All-Star portion** of each season from 2000 to 2024.
- Includes **traditional box score stats** (PTS, REB, AST, etc.), **advanced metrics** (PIE, TS%, Net Rating), **team stats** (W-L, Team TO%, Pace, etc.), and **player demographics** (Age, Team, Season, Player ID).

We used one-tailed t-tests to confirm which metrics were relevant to All-Star selection, paired t-tests to determine consistency throughout seasons, and chi-square tests to confirm the impact of high average scores for selection. Leveraging the same feature set, we fit a regularized linear regression model (Ridge and Lasso) to quantify how much of All-Star selection can be explained by box-score stats, and we train a k-Nearest Neighbors classifier to predict All-Star status with cross-validated accuracy.

Analysis

First, we ran **one-tailed t-tests** on the average of significant player features across seasons. Our null hypothesis stated that the average chosen feature (TOTAL_MIN, PTS, etc) for All-Stars is less than or equal to that of non-All-Stars, while the alternative hypothesis held that All-Stars achieve higher season statistics. *Figure 1* shows the resulting t-stat and p-values after running these hypotheses. Observe that for the chosen features **PTS** and **USG_PCT** we achieved the highest t-stats (63.765 and 43.431 respectively) with p-values converging to 0, leading us to reject the null hypothesis. We conclude that **All-Star players average higher points per game and usage percentage than non-All-Star players**.

	t-stat	p-value
TOTAL_MIN	40.264	0.00E+00
PTS	63.765	0.00E+00
PIE	21.676	2.19E-102
NET_RATING	13.137	1.98E-39
WINS_EST	28.629	2.77E-174
USG_PCT	43.431	0.00E+00
FTA	68.06	0.00E+00
REB	38.522	3.89E-305
TS_PCT	12.534	4.26E-36
TS_minus_EFG	10.346	2.82E-25

Figure 1. One-tailed t-test results of each statistic against All-Stars

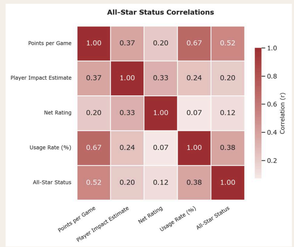

Figure 2. Heatmap showing PCC correlation between All-Star status and key features

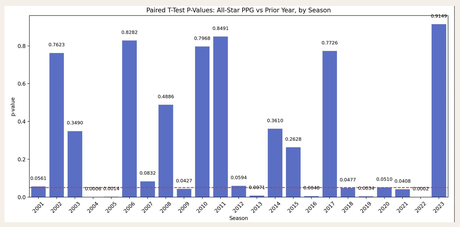

Figure 3. Bar chart of p-values from paired t-test of All-Star PPG vs Prior Year (2001-2023)

We then created a **quadrant analysis visualization** using a **scatterplot** of players' usage percentage against net rating (*Figure 4*). Our goal here was to test the **null hypothesis** that winning does not have an outsized impact on All-Star selection for high usage players, against the **alternative** that a combination of high usage and impact on winning is strongly correlated with All-Star selection. We filtered players with ≥200 total minutes, then divided them into four quadrants based on average usage rate and positive/negative net rating. We highlighted All-Star players in red, while non-All-Stars appear in light gray, with point sizes corresponding to scoring output. The visualization clearly shows that players with **high usage on winning teams are far more likely to be selected as All-Stars (24.8%)** when compared with the other three quadrants, including the "**High Usage, Losing Teams**" quadrant **(3.6%)**.

Figure 1 shows that most of the chosen features exhibit high t-stats and low p-values. We used the **Pearson correlation coefficient** to measure correlation between high-offense metrics and All-Star status, shown in *Figure 2*, and found that out of the chosen features in *Figure 1*, **PTS** (Points per game), **USG_PCT** (usage rate %), **PIE** (Player Impact Estimate), and **NET_RATING** (net rating) have the highest r-values with All-Star status in decreasing order. Possible explanations for TOTAL_MIN, FTA, and other features having high t-stats is imbalanced team performance and foul preference towards star players.

Next, we examined whether All-Star players' scoring changed from the previous season to this one by using a **paired t-test** on each player's PPG difference. Our null hypothesis asserted that the average change in PPG (this season minus last season) is zero, while the alternative allowed for any non-zero change. *Figure 3* shows us that there were only 7 years where our null hypothesis can be rejected, whereas in most years **there's no statistically significant evidence that All-Star players' scoring this season differs from their scoring last season.**

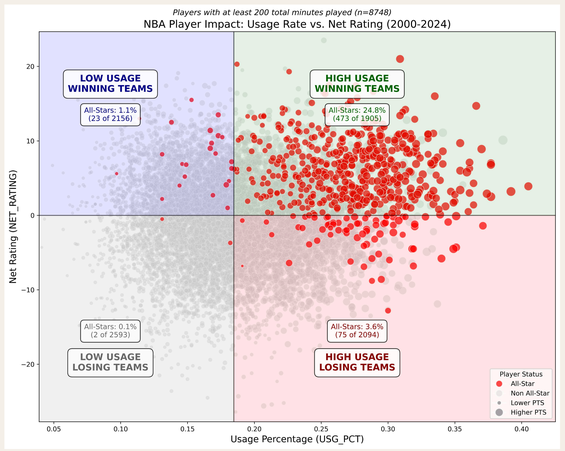

Figure 4. Scatterplot of Usage Percentage vs Net Rating. Red dots are All-Stars. Bigger dots indicate higher Points Per Game.

Model Analysis

Model	Fold 1 R^2	Fold 2 R^2	Fold 3 R^2	Fold 4 R^2	Fold 5 R^2	Mean R^2
Linear Regression	0.246	0.202	0.317	0.263	0.286	0.263
Ridge Regression (L2)	0.246	0.202	0.317	0.263	0.286	0.263
Lasso Regression (L1)	0.243	0.212	0.293	0.258	0.276	0.256
Polynomial + Ridge Regression	0.471	0.443	0.532	0.49	0.516	0.491

Figure 5. R^2 values under 5-fold CV for OLS, Ridge, Lasso, and Polynomial + Ridge Regression

Next, we performed an **Ordinary Least Squares regression** using our four key features to establish a baseline. Using **5-fold cross-validation**, the model achieved a mean R^2 of 0.263 (*Figure 5*). To capture more complex interactions, we extended the model with **L2 polynomial** features and **Ridge regression**, ultimately finding that coupling the two gave us the best results. This nearly doubled performance, raising the mean R^2 to 0.49 and yielding a more confident explanation of All-Star selection variance, following from the highly non-linear nature of the relationships.

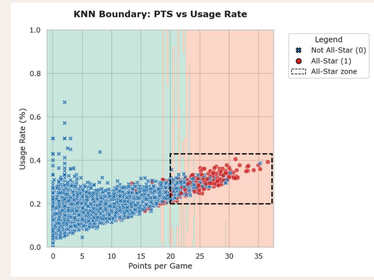

Figure 6. KNN decision-boundary plot labeled with PTS and USG_PCT

Finally, we ran the **K-Nearest Neighbors** algorithm with k=5. In cross-validation, this classifier correctly identified players' All-Star status with a mean accuracy of 0.967. Achieving a near-perfect accuracy with only 4 features demonstrates that **these high-offense metrics alone form an exceptionally strong predictive signal for All-Star status**. *Figure 6* clearly shows an "All-Star Zone" where PTS and USG_PCT start to reflect All-Star status.

Limitations

- **Binary nature** of the our classification makes ML models limited in what we can test (regression lines of best fit are bad for visualizations, ML)
- The correlation between our metrics and All-Star selection doesn't necessarily indicate causation, as **other unmeasured factors** (media exposure, market size, career reputation) likely influence selection
- A key limitation is **multicollinearity** among some of our variables (Usage Rate inherently contains Minutes, FGA, and FTA), complicating efforts to **isolate their independent effects** on All-Star selection.
- Another limitation was the restricted access to comprehensive NBA statistics. While the NBA Official API provided a solid foundation, many alternative APIs required paid access or lacked sufficient historical data across the full range of seasons we analyzed. This constrained the depth and variety of features we could include in our models.

Conclusion

- Our analysis shows that **high-offensive involvement metrics and positively impacting the team like points per game, usage rate, and net rating are strong predictors of NBA All-Star selection**. Both our hypothesis tests and machine learning models, especially **Ridge regression and KNN**, consistently highlighted these stats as the most influential.
- That said, we acknowledge that stats alone don't tell the full story. Factors like **fan voting, media exposure, and market size** likely play important roles that aren't captured in our models. Players with strong off-court impact or narrative appeal may be overlooked if only performance data is considered.
- Moving forward, we see value in incorporating additional features such as social media engagement or market visibility, and in testing **non-linear models to uncover more complex patterns**. Our work provides a strong statistical foundation and opens the door to understanding the broader forces that shape All-Star recognition